

CSC411: Project #2

Due on Friday, March 10, 2017

Fiona Leung *999766061*

Timur Borkhodoev *999376394*


March 10, 2017

Part 1

Dataset description The dataset contained an even number of images per digit, having no more or less information the neural networks could train off of to determine one digit better apart from another. Each digit had multiple images of it where it was drawn with varying writing styles. Some digits were written with more loops and imperfections than others, had different thicknesses or were drawn slanted at different angles. Some digit images had discontinuous lines in the number shapes.

For example, in the images of the 0's some had a closed loop while others were perfect and had a small edge that jutted out near its top.

1. variety of angles
2. different styles of handwriting
3. gaps between continuous lines
4. different thickness levels



Part 2

Compute the network function by propagating forward and discarding intermediate results

```

1 def compute_network(x, W0, b0, W1, b1):
2     _,_, output = forward(x, W0, b0, W1, b1)
3     return argmax(output)

```

Part 3

1. We will use negative log-probabilities as our cost function, and find its gradient

```

1  def cross_entropy(y, y_):
2      return -sum(y_ * log(y))

```

2. Vectorized code for computing gradient of the cost function

```

1  def deriv_multilayer(W0, b0, W1, b1, x, L0, L1, y, y_):
2      dCdL1 = y - y_
3      dCdW1 = dot(L0, dCdL1.T)
4      dCdobydodh = dot(W1, dCdL1)
5      diff = 1 - L0**2
6
7      dCdW0 = tile(dCdobydodh, 28 * 28).T * dot(x, (diff.T))
8      dCdb1 = dCdL1
9      dCdb0 = dCdobydodh * diff
10
11     return dCdW1, dCdb1, dCdW0, dCdb0

```

Part 4

```

1  def train(plot=False):
2      global W0, b0, W1, b1
3      global plot_iters, plot_performance
4      plot_iters = []
5      plot_performance = []
6      alpha = 1e-3
7      for i in range(150):
8          X, Y, examples_n = get_batch(i * 5, 10)
9
10         update = np.zeros(4)
11
12         for j in range(examples_n):
13             y = Y[j].reshape((10, 1))
14             x = X[j].reshape((28 * 28, 1)) / 255.
15             L0, L1, output = forward(x, W0, b0, W1, b1)
16             gradients = deriv_multilayer(W0, b0, W1, b1, x, L0, L1, output, y)
17             update = [update[k] + gradients[k] for k in range(len(gradients))]
18
19         # update the weights
20         W1 -= alpha * update[0]

```

```


21         b1 -= alpha * update[1]
22         W0 -= alpha * update[2]
23         b0 -= alpha * update[3]
24         if plot:
25             plot_iters.append(i * examples_n)
26             plot_performance.append(test_perf())
27     return plot_iters, plot_performance
28 train(plot=True)

```

```

1 def get_batch(offset, example_per_class=5):
2     # 5 examples per class
3     classes_num = 10
4     x_batch = np.zeros((example_per_class * classes_num, 28 * 28))
5     y_batch = np.zeros((example_per_class * classes_num, classes_num))
6     for i in range(classes_num):
7         for j in range(example_per_class):
8             x_batch[i * example_per_class + j] = M['train' + str(i)][j +
9                                                         offset]
10            y_batch[i * example_per_class + j][i] = 1
11     return x_batch, y_batch, example_per_class * classes_num

```



Part 5

So as discussed in lecture large errors are penalized quadratically in linear regressions. So our multinomial regression doesn't suffer from it. We start with generating a noise for our data set.

```
1  # generate noise and  $N(0, \sigma^2)$ 
2  noise = scipy.stats.norm.rvs(scale=5, size=784*50*10)
3  noise = noise.reshape(500, 784)
4  X, Y, n = get_batch(offset=0, examples=50)
5  X += noise
```

After modifying our data set with noise - we get images that look like this

